

Agentic AI for Lead Generation

Partner companies for the Marketing Chair

Tobias Weiß · Guest Lecture · Marketing Chair

45 minutes · BWL Master students

Why this session exists

- The chair partners with companies: funding, data, cases
- Finding + qualifying partners is manual, slow, unscalable
- Agentic AI can research, score and rank prospects
- **Demo output = input for your Bachelor seminar**

Nobody wins from another lecture about chatbots.

What you will be able to do

By the end...	You can...
Understand	Explain what "agentic" means, and why it is not a chatbot
Operate	Read how a harness (pi) gives an AI tools, rules and a task
Assess	Judge an agentic output: checked? citable? repeatable?
Apply	Reuse the loop from the demo (contract → run → verify) for partner sourcing


Model = brain. Harness = hands. Workflow = muscles.

Roadmap — 35 min + 10 min Q&A

Block	Content	Time
1 · Foundations	agentic ≠ chat · model · harness · loop	~10 min
2 · Method	contracts · the lead-gen loop · evidence	~7 min
3 · Walkthrough	live demo: contract → agent → verify	~13 min
4 · Reflection	providers · take-aways · resources	~5 min
5 · Q&A	your questions · discussion	10 min


The walkthrough is the heart — theory is fuel for it. **10 minutes are reserved for your questions.**

Agentic ≠ chat



"Agentic" is a working environment — not a better chat.

The parts: model · harness · loop



Harnesses (2026): pi, OpenCode, open-source ecosystems — same models, different results, depending on the harness.


Why now — the evidence inside the research corpus

13,096 marketing papers · 7,771 published in 2026 alone

Signal	Count
Papers mentioning agentic AI	91
Papers on lead generation	27
Both (agentic × lead-gen)	0

The field is exploding — and nobody connects the two yet. This chair can be first.

How to talk to the machine: prompts vs contracts




Prompt = how you ask a chat · Contract = how you run an agent.

A lead-generation loop as an agentic workflow

LEAD-GENERATION LOOP

agent executes steps 2–5 · the human owns steps 1 and 6



Six steps — the agent executes 2–5, you own definition and evaluation.

Now: a live walkthrough

Task: build the **partner-candidate shortlist** (top 10 B2B companies) with:

- Fit signal per company vs. the chair's partnership goals
- Named source URL per claim (citable, verifiable)
- Ranked output — your input for the Bachelor seminar

I run this live in a harness — what you see is what the agent does.

Walkthrough · Step 1 — The contract

AGENTS.md / brief (what the agent reads at start)

Goal: shortlist of 10 B2B partner candidates for the chair

Criteria: industry fit, innovation agenda, university links

Sources: company sites, press, funding registries

Output: ranked list, every claim with a URL

Rules: no invented facts; if unsure, mark as unverified

Done means: verified, cited, repeatable.

Walkthrough · Step 2 — The agent at work

Watch the six-step loop run live:

- **only watch** — no magic, every tick is visible in the terminal
- agent **opens the browser** (MCP), reads real company pages
- scores against **the contract** → markdown table with **URLs**
- **iterates**: pause, correct, re-check — that is the loop

Nothing happens off-screen — the loop is fully visible.

Walkthrough · Step 3 — The three questions

1. **Checked?** Did the agent verify its output, not just generate it?
2. **Citable?** Does every claim have a URL that resolves?
3. **Repeatable?** Re-run the same contract → same result?

The final slide of the walkthrough: apply them to the live result.

What you just saw

1. A contract (rules) — not a prompt
2. An agent doing the loop: plan → act → check → correct
3. A verifiable artefact — the ranked shortlist

That is the method — reusable for any partner-sourcing task.

What the chair gets out of it

- A reproducible partner-sourcing method (not a one-off search)
- A ranked, sourced shortlist — usable immediately
- Students who can run agentic workflows (a real skill for the CV)
- Input data for the **Bachelor seminar** — thematically linked

The lecture is not the output. The shortlist is.

Key AI providers — what you can use as a student

Provider / offer	What you get	Where
JLU HRZ API-Service ★	LLMs for all JLU members: free local Qwen models (256K–1M ctx) + commercial, works with OpenCode / Claude Code	uni-giessen.de HRZ API-Service · access via ki@uni-giessen.de
Google AI student plan	1 year Google AI Plus free · Gemini, higher limits, 400 GB	one.google.com/ai-student — verify via SheerID (student status)
OpenRouter	one API key → many models (OpenAI, Anthropic, Google, open) · free models + pay-as-you-go	openrouter.ai
Harnesses	the tools from today's demo — OpenCode , pi (work with any key above)	opencode.ai · pi.dev
Local models (Ollama)	free, run on your own laptop, no account, private	ollama.com

Start with **HRZ API**: free for JLU members · Google deal expires **Dec 31, 2026**.

Take-aways

1. Agentic = working environment, not a chatbot
2. Harness = model + tools + rules = accountable AI work
3. Contract first: define "done" before you start
4. Verify, cite, repeat — the three questions

Next week: this shortlist becomes the Bachelor seminar's input.

Resources

- **JLU HRZ API-Service** (free LLMs for JLU members) — uni-giessen.de → HRZ → KI → API-Service · model overview
- **pi** — pi.dev · docs: pi.dev/docs · **OpenCode** — opencode.ai · **OpenRouter** (one key, many models) — openrouter.ai
- **Google AI student plan** — one.google.com/ai-student (1 year free for students)
- **marketing-research** (chair corpus, public) — github.com/tobias-weiss-ai-xr/marketing-research
- **skeleton-research** (forkable corpus skeleton) — github.com/tobias-weiss-ai-xr/skeleton-research
- **This deck** — `lecture/` in this corpus repo (English, 45 min)

Thank you!

Questions?

Tobias Weiß · guest lecture · Marketing Chair

Agentic AI for Lead Generation · 45 min